

## **...And Optimize**

---

'Cause we are not just dealing with ML, ain't we?



# Our Current Situation

**The results so far are not comforting**

...But it's worth seeing what is going on **over time**




# Our Current Situation

The results so far are not comforting

...But it's worth seeing what is going on over time

```
In [2]: stop = 1400  
util.plot_rul(tr_pred[:stop], tr['rul'][:stop], figsize=figsize)
```



✎ ..And we get the same shapes also on the validation and test set

## ...And Two Observations about It

- All sequences of predicted RULs start with a "flat" section
- ...And after a while they bend and start decreasing

**First: why is this happening? And why is it so consistent?**



## ...And Two Observations about It

- All sequences of predicted RULs start with a "flat" section
- ...And after a while they bend and start decreasing

**First: why is this happening? And why is it so consistent?**

**One reason is that large RUL values are under-represented**

...Since not all machines run for the same time

- As a result we will have larger noise
- ...And it may be impossible to predict RULs larger than in the training set

However, the curve bend relatively late

- ...And therefore there must be something more



## ...And Two Observations about It

- All sequences of predicted RULs start with a "flat" section
- ...And after a while they bend and start decreasing

**First: why is this happening? And why is it so consistent?**



## ...And Two Observations about It

- All sequences of predicted RULs start with a "flat" section
- ...And after a while they bend and start decreasing

**First: why is this happening? And why is it so consistent?**

**The main reason is that degradation does not start immediately**

...But typically only when microscopic defects grow to become perceivable

- As a result, early on the NN will be "see" examples with comparable input
- ...But different target values

When an MSE loss, the optimal choice in this case is to predict the average



## ...And Two Observations about It

- All sequences of predicted RULs start with a "flat" section
- ...And after a while they bend and start decreasing

**Second: is this pattern good or bad news for us?**

**Our goal is **not** to regress RUL values with high accuracy**

...But rather to define a maintenance **policy** in the form:

$$f(x, \omega) < \theta \Rightarrow \text{trigger maintenance}$$

- For this, we just need to **stop at the right time**
- ...And our model may be accurate enough in the region that matters





# Threshold Calibration as an Optimization Problem

## Given a RUL estimator

...We can choose when to trigger maintenance by calibrating  $\theta$

- This is in fact an (other) optimization problem
- ...And to formulate it we need a cost function

## Our cost function will rely on this simplified cost model:

- Whenever a turbine operates for a time step, we gain a profit of 1 unit
- A failure costs  $C$  units (i.e. the equivalent of  $C$  operation days)
- We never trigger maintenance before  $s$  time steps

Some comments:

- $C$  is actually an offset over the cost of maintenance
- The last rule mimics using preventive maintenance as a fail-safe mechanism



# The Cost Function

Let  $x_k$  be the times series for machine  $k$  (out of  $n_r$ ), and  $n_t$  its length

With our RUL based policy:

- Given a cost function  $cost(f(x_k), x_k, \theta)$  for one machine, the total cost is:

$$\sum_{k=1}^{n_r} cost(f(x_k), x_k, \theta)$$

- The time step when we trigger maintenance is given by:

$$\min\{i = 1..n_t \mid f(x_{ki}) < \theta\}$$

- A failure occurs if:

$$f(x_{ki}) \geq \theta \quad \forall i = 1..n_t$$



# The Cost Function

The cost formula **for a single machine** will be:

$$cost(f(x_k), x_k, \theta) = op\_profit(f(x_k), \theta) + fail\_cost(f(x_k), \theta)$$

Where:

$$op\_profit(f(x_k), \theta) = -\max(0, \min\{i \in I_k \mid f(x_{ki}) < \theta\} - s)$$

$$fail\_cost(f(x_k), \theta) = \begin{cases} C & \text{if } f(x_{ki}) \geq \theta \quad \forall i \in I_k \\ 0 & \text{otherwise} \end{cases}$$

- $s$  units of machine operation are guaranteed
- ...So we gain over the default policy only if we stop after that
- Profit is modeled as a negative cost



# The Cost Function

**Normally, we would determine  $s$  and  $C$  by talking to a domain expert**

...In our case we will pick reasonable values based on our data

- First, we collect all failure times:

```
In [3]: tr_failtimes = tr.groupby('machine')['cycle'].max()
```

- Then, we define  $s$  and  $C$  based on statistics:

```
In [4]: safe_interval = tr_failtimes.min()  
maintenance_cost = tr_failtimes.max()
```

- For the safe interval  $s$ , we choose the minimum failure time
- For the maintenance cost  $C$  we choose the largest failure time

We are talking about jet engines, so failing is BAD



# Calibration and Policy Definition Problem

Our calibration problem is then in the form:

$$\operatorname{argmin}_{\theta} \sum_{k=1}^{n_r} \operatorname{cost}(f(x_k), x_k, \theta)$$

- If we pair it with our previous training step
- ...We obtain a formulation for the entire **policy definition** problem:

$$\operatorname{argmin}_{\theta} \sum_{k=1}^{n_r} \operatorname{cost}(f(x_k, \omega^*), x_k, \theta)$$

where:  $\omega^* = \operatorname{argmin}_{\omega} \{ L(y, \hat{y}) \mid y = f(x, \omega) \}$

This is how we should have started in the first place



# Solving the Calibration Problem

Solving the calibration problem is very easy:

$$\operatorname{argmin}_{\theta} \sum_{k=1}^{n_r} \operatorname{cost}(f(x_k, \omega^*), x_k, \theta)$$

where:  $\omega^* = \operatorname{argmin}_{\omega} \{ L(y, \hat{y}) \mid y = f(x, \omega) \}$

- We need to optimize a single (scalar) variable, i.e.  $\theta$
- ...And changing  $\theta$  does not impact the optimal  $\omega$

**This is a **univariate optimization** problem**

- The cost function is non-differentiable
- ...But the problem is so simple that even **grid search** will work very well



# Solving the Calibration Problem


We can **sample a range** of values for the  $\theta$  parameter

...Then simply pick the value with the smallest cost

■ The code in `optimize_threshold` can also plot the corresponding cost surface

```
In [6]: cmodel = util.RULCostModel(maintenance_cost=maintenance_cost, safe_interval=safe_interval)
th_range = np.linspace(-2, 40, 100)
tr_thr = util.optimize_threshold(tr['machine'].values, tr_pred, th_range, cmodel, plot=True, fig=fig)
print(f'Optimal threshold for the training set: {tr_thr:.2f}')
```

Optimal threshold for the training set: 3.52



# Evaluation

Finally, we can check how we are doing on the test set:

```
In [7]: tr_c, tr_f, tr_sl = cmodel.cost(tr['machine'].values, tr_pred, tr_thr, return_margin=True)
        ts_c, ts_f, ts_sl = cmodel.cost(ts['machine'].values, ts_pred, tr_thr, return_margin=True)
        print(f'Cost: {tr_c} (training), {ts_c} (test)')
```

```
Cost: -18238 (training), -7075 (test)
```

We can also evaluate the margin for improvement:

```
In [8]: print(f'Avg. fails: {tr_f/len(tr_mcn):.2f} (training), {ts_f/len(ts_mcn):.2f} (test)')
        print(f'Avg. slack: {tr_sl/len(tr_mcn):.2f} (training), {ts_sl/len(ts_mcn):.2f} (test)')
```

```
Avg. fails: 0.01 (training), 0.00 (test)
Avg. slack: 15.06 (training), 11.63 (test)
```

- Slack = distance between when we stop and the failure
- The results are actually quite good!

 ... And we also generalize fairly well



## Some Considerations

**In principle, RUL regression is a very hard problem**

- Our linearly decreasing RUL assumption is just a rough oversimplification
- ...RUL is inherently subject to stochasticity
- ...And depends on the how the machine **will be** used

**But we **don't care**, since RUL prediction was **not our true problem****

The real problem involved both **prediction and optimization**

- We had to optimize the NN parameters (to obtain good predictions)
- We had to optimize the threshold

The ultimate goal was to **reduce maintenance cost**

**Keep in mind **the big picture****

- In a "predict, then optimize" setting
- ...Quality should be judged on the final cost

